



Code Along - Machine Learning on AWS

By Parag Pansare

Data Science Professional & AWS Cloud Architect

Agenda

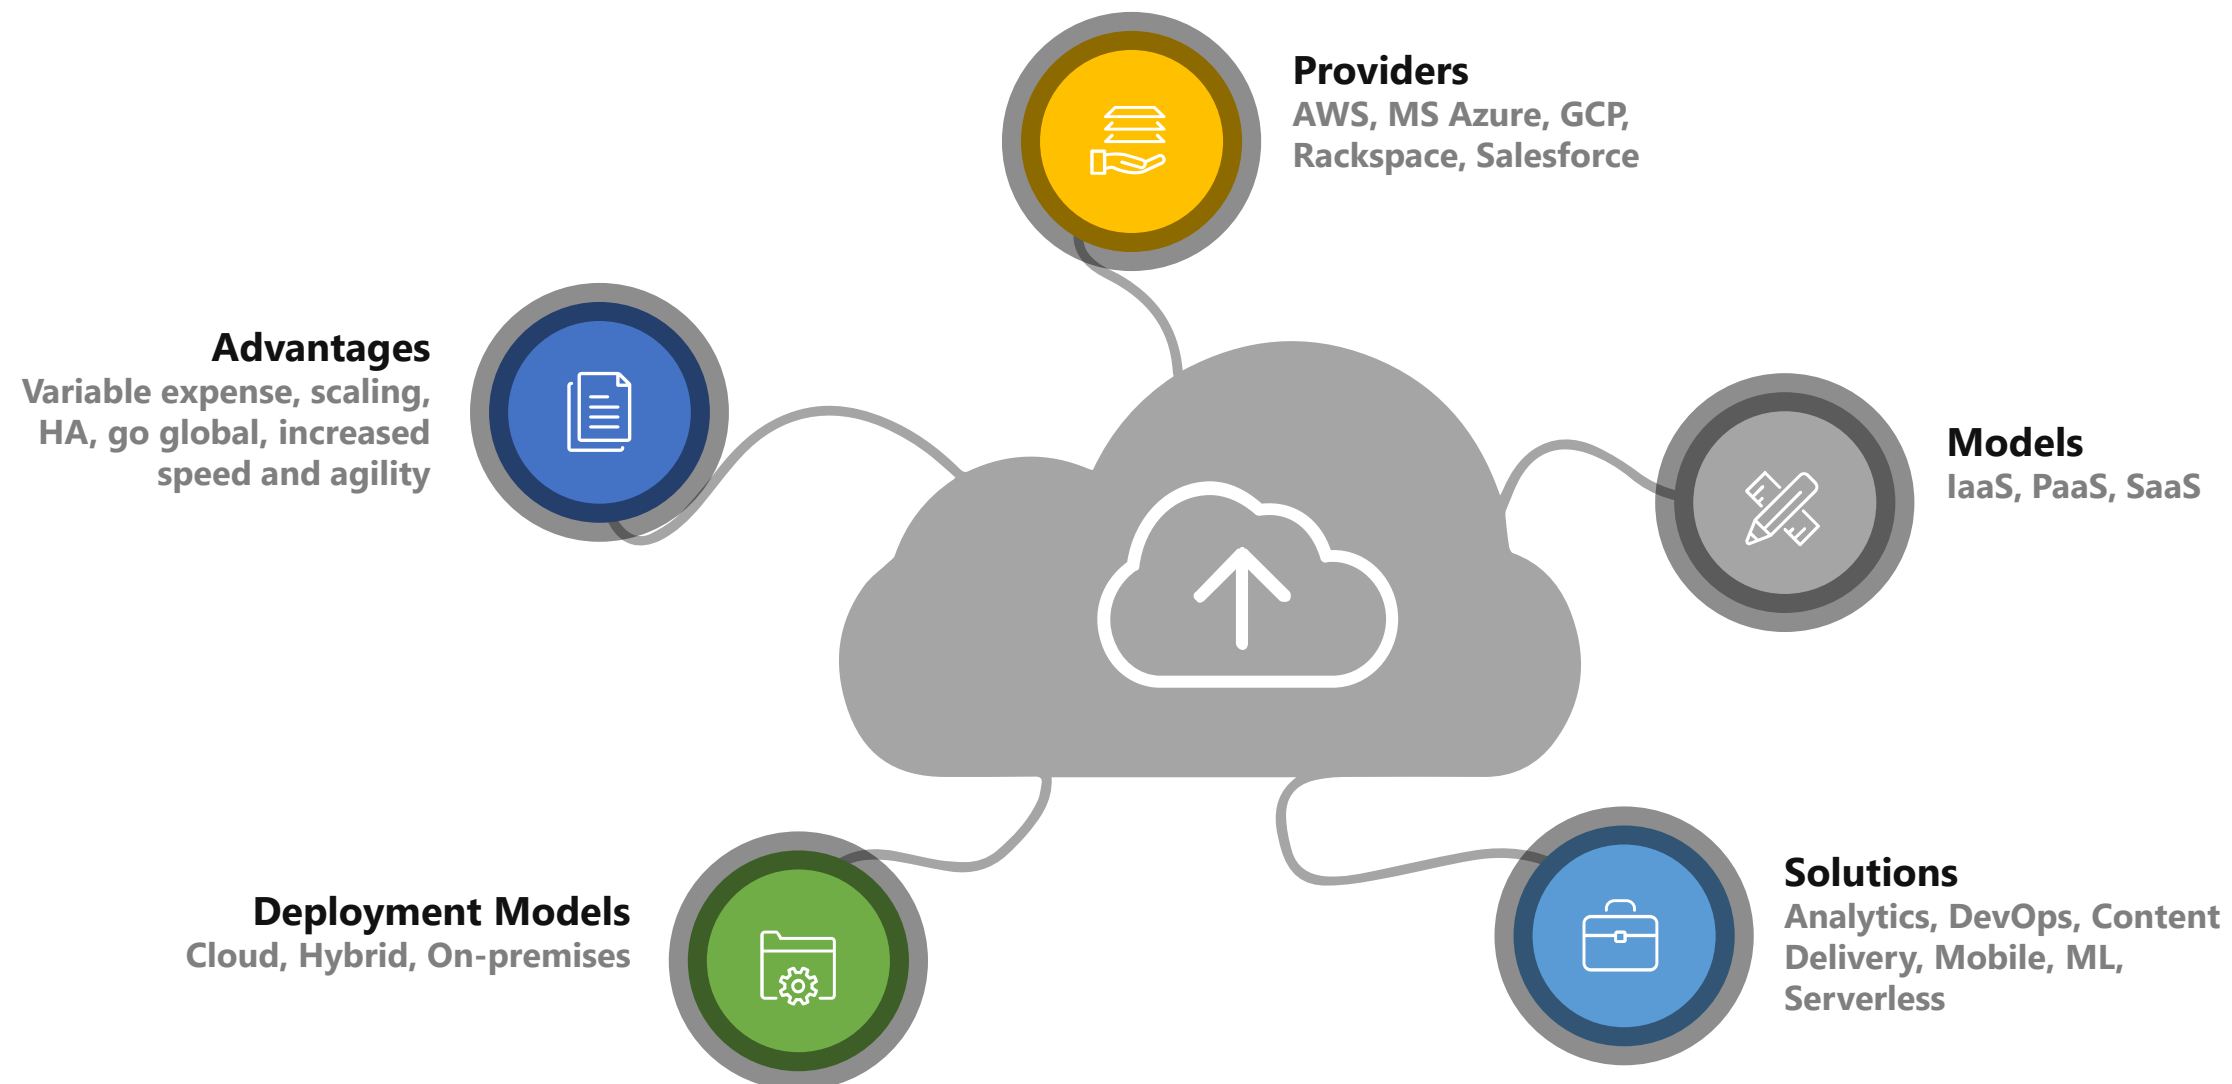
- Overview:
 - What is Cloud Computing and AWS?
 - What is Machine Learning?
- Machine Learning using AWS on Amazon SageMaker:
 - Overview and Architecture
 - Hands on problem solving using ML on Python
 - Programming Model
 - K-Means Algorithm
 - Automatic Model Tuning/Hyperparameter Tuning
- Machine Learning using other AWS services

Prerequisite

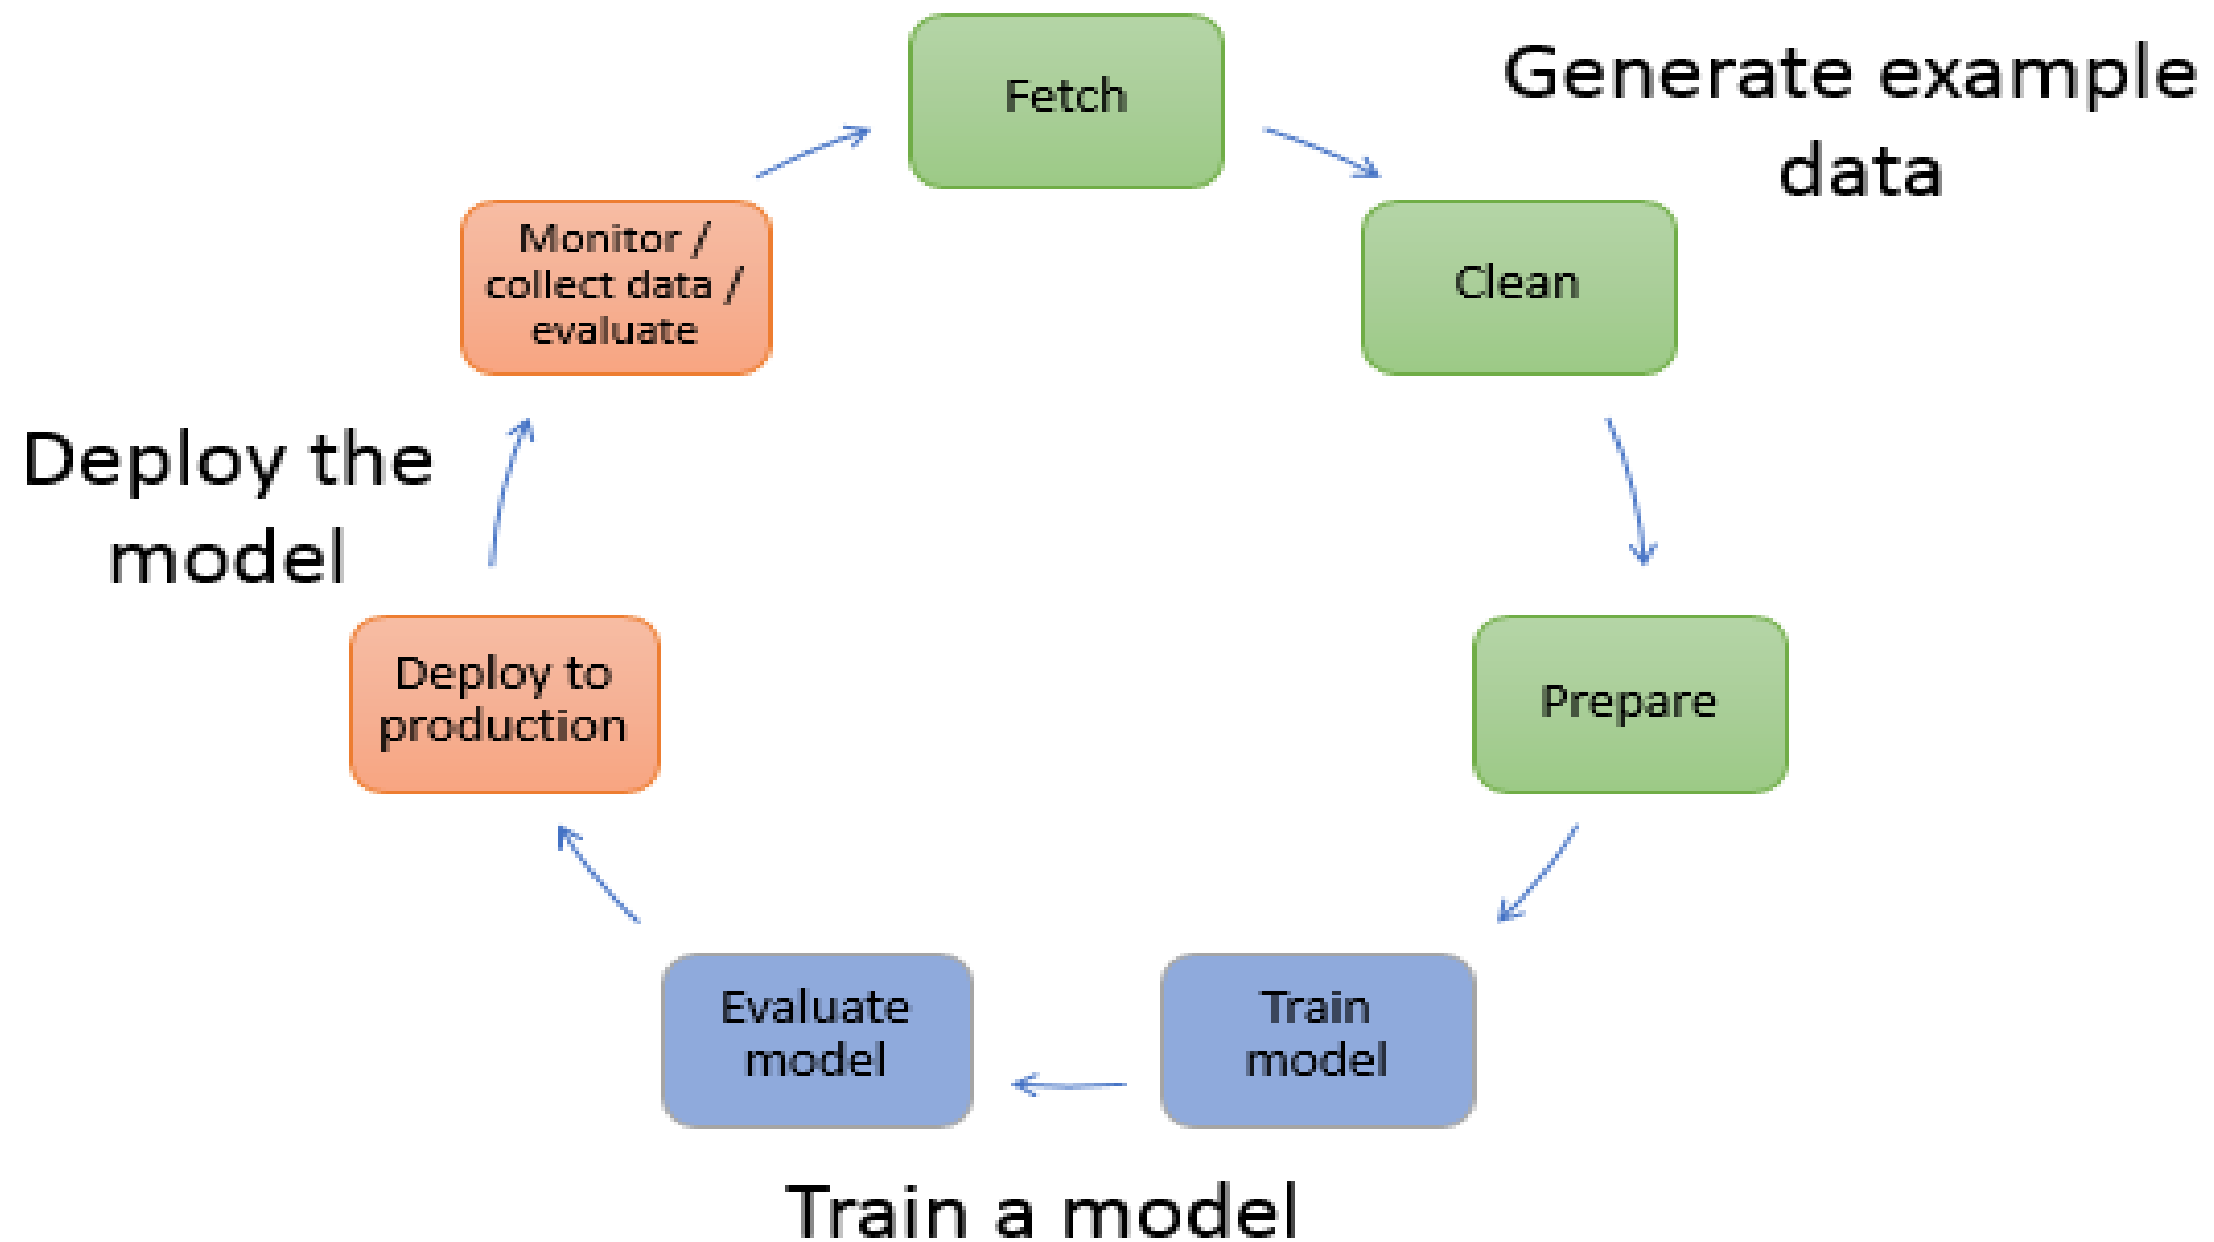
- AWS account
- Basic understanding of Python and Jupyter Notebook
- Basic understanding of Machine Learning and AWS
- Healthy lunch ;)

What is Cloud Computing and AWS?

On-demand delivery of compute power, database storage, apps, and other IT resources through a cloud services platform via internet with pay-as-you-go pricing

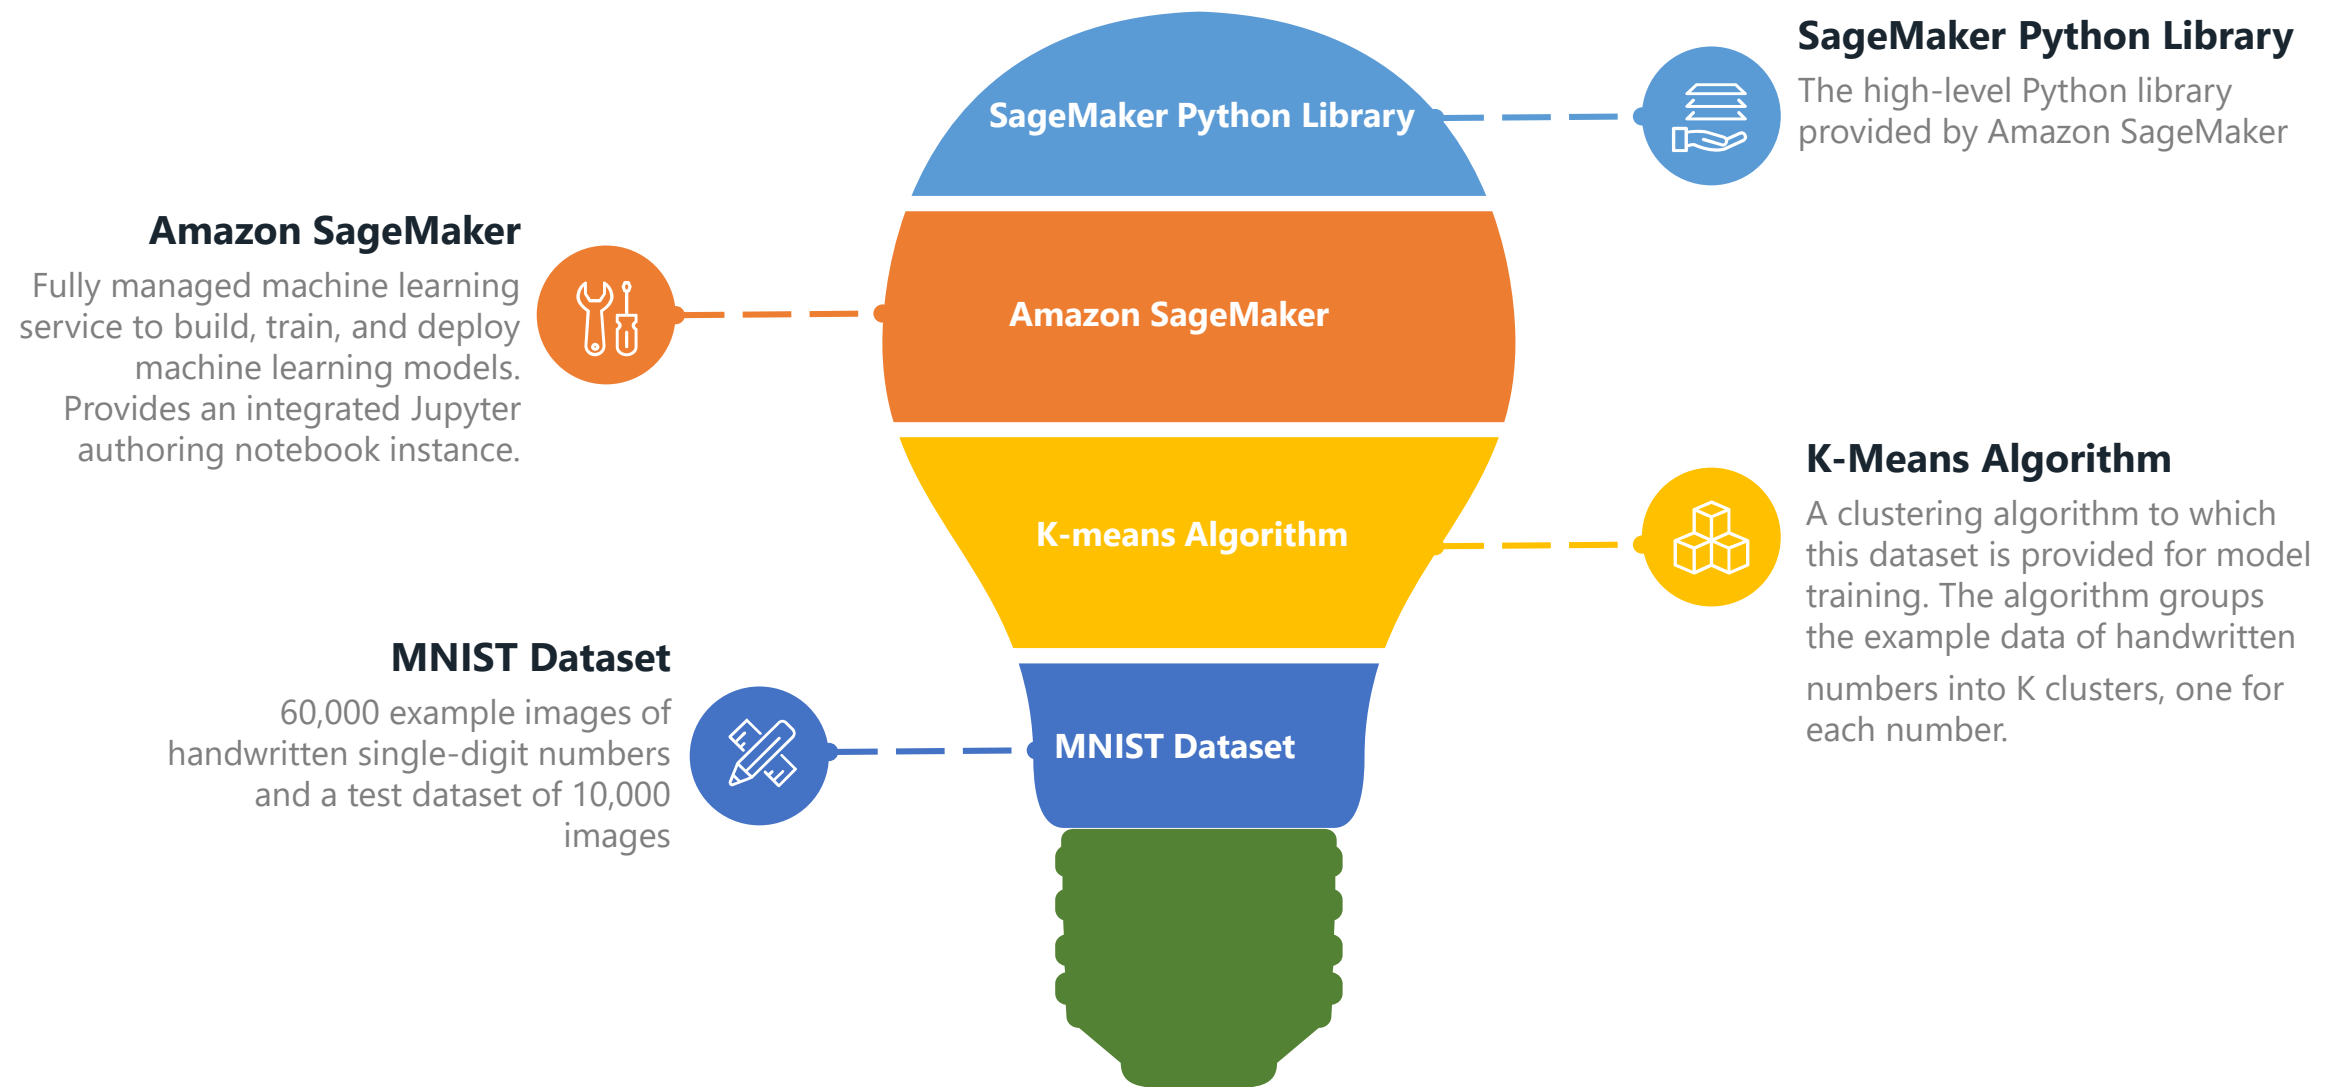


What is Machine Learning?



Problem Statement...

Classification problem to predict the number of a handwritten digit image



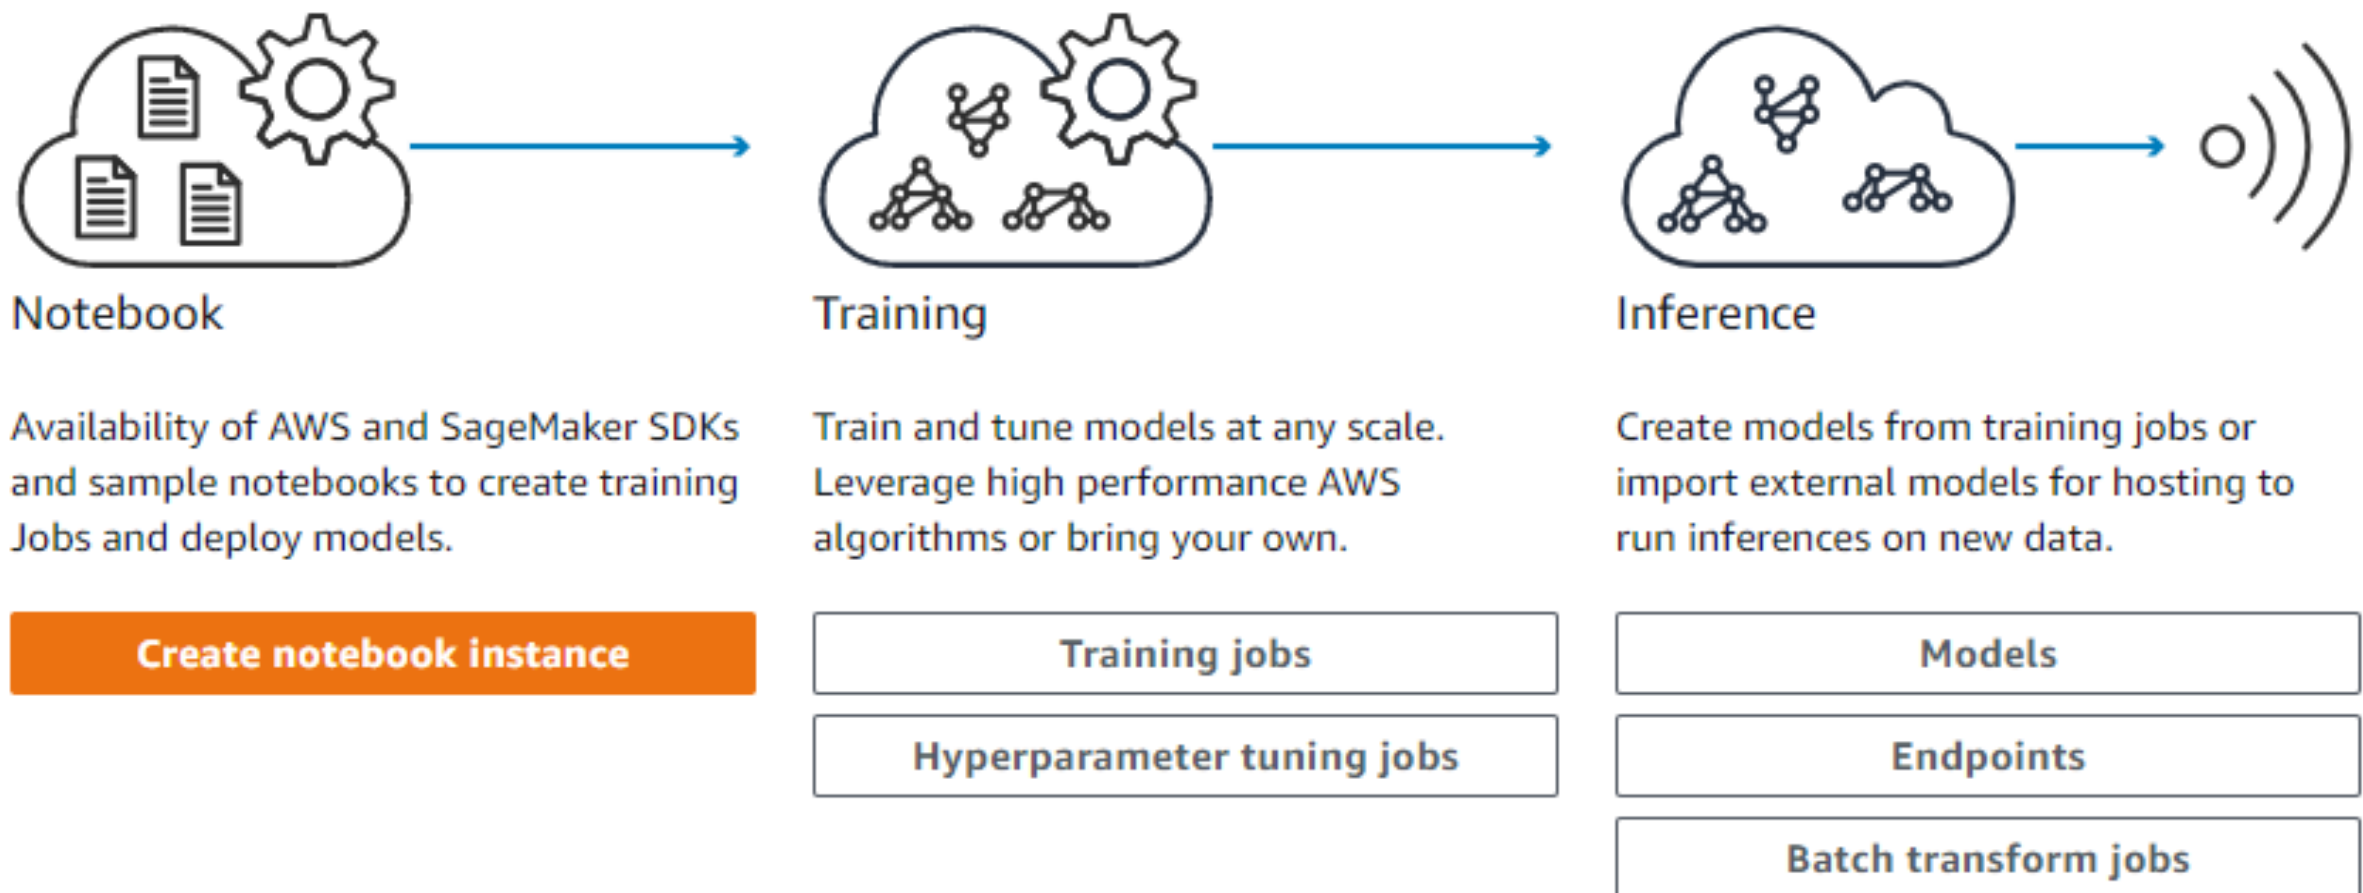


Machine Learning using Amazon SageMaker

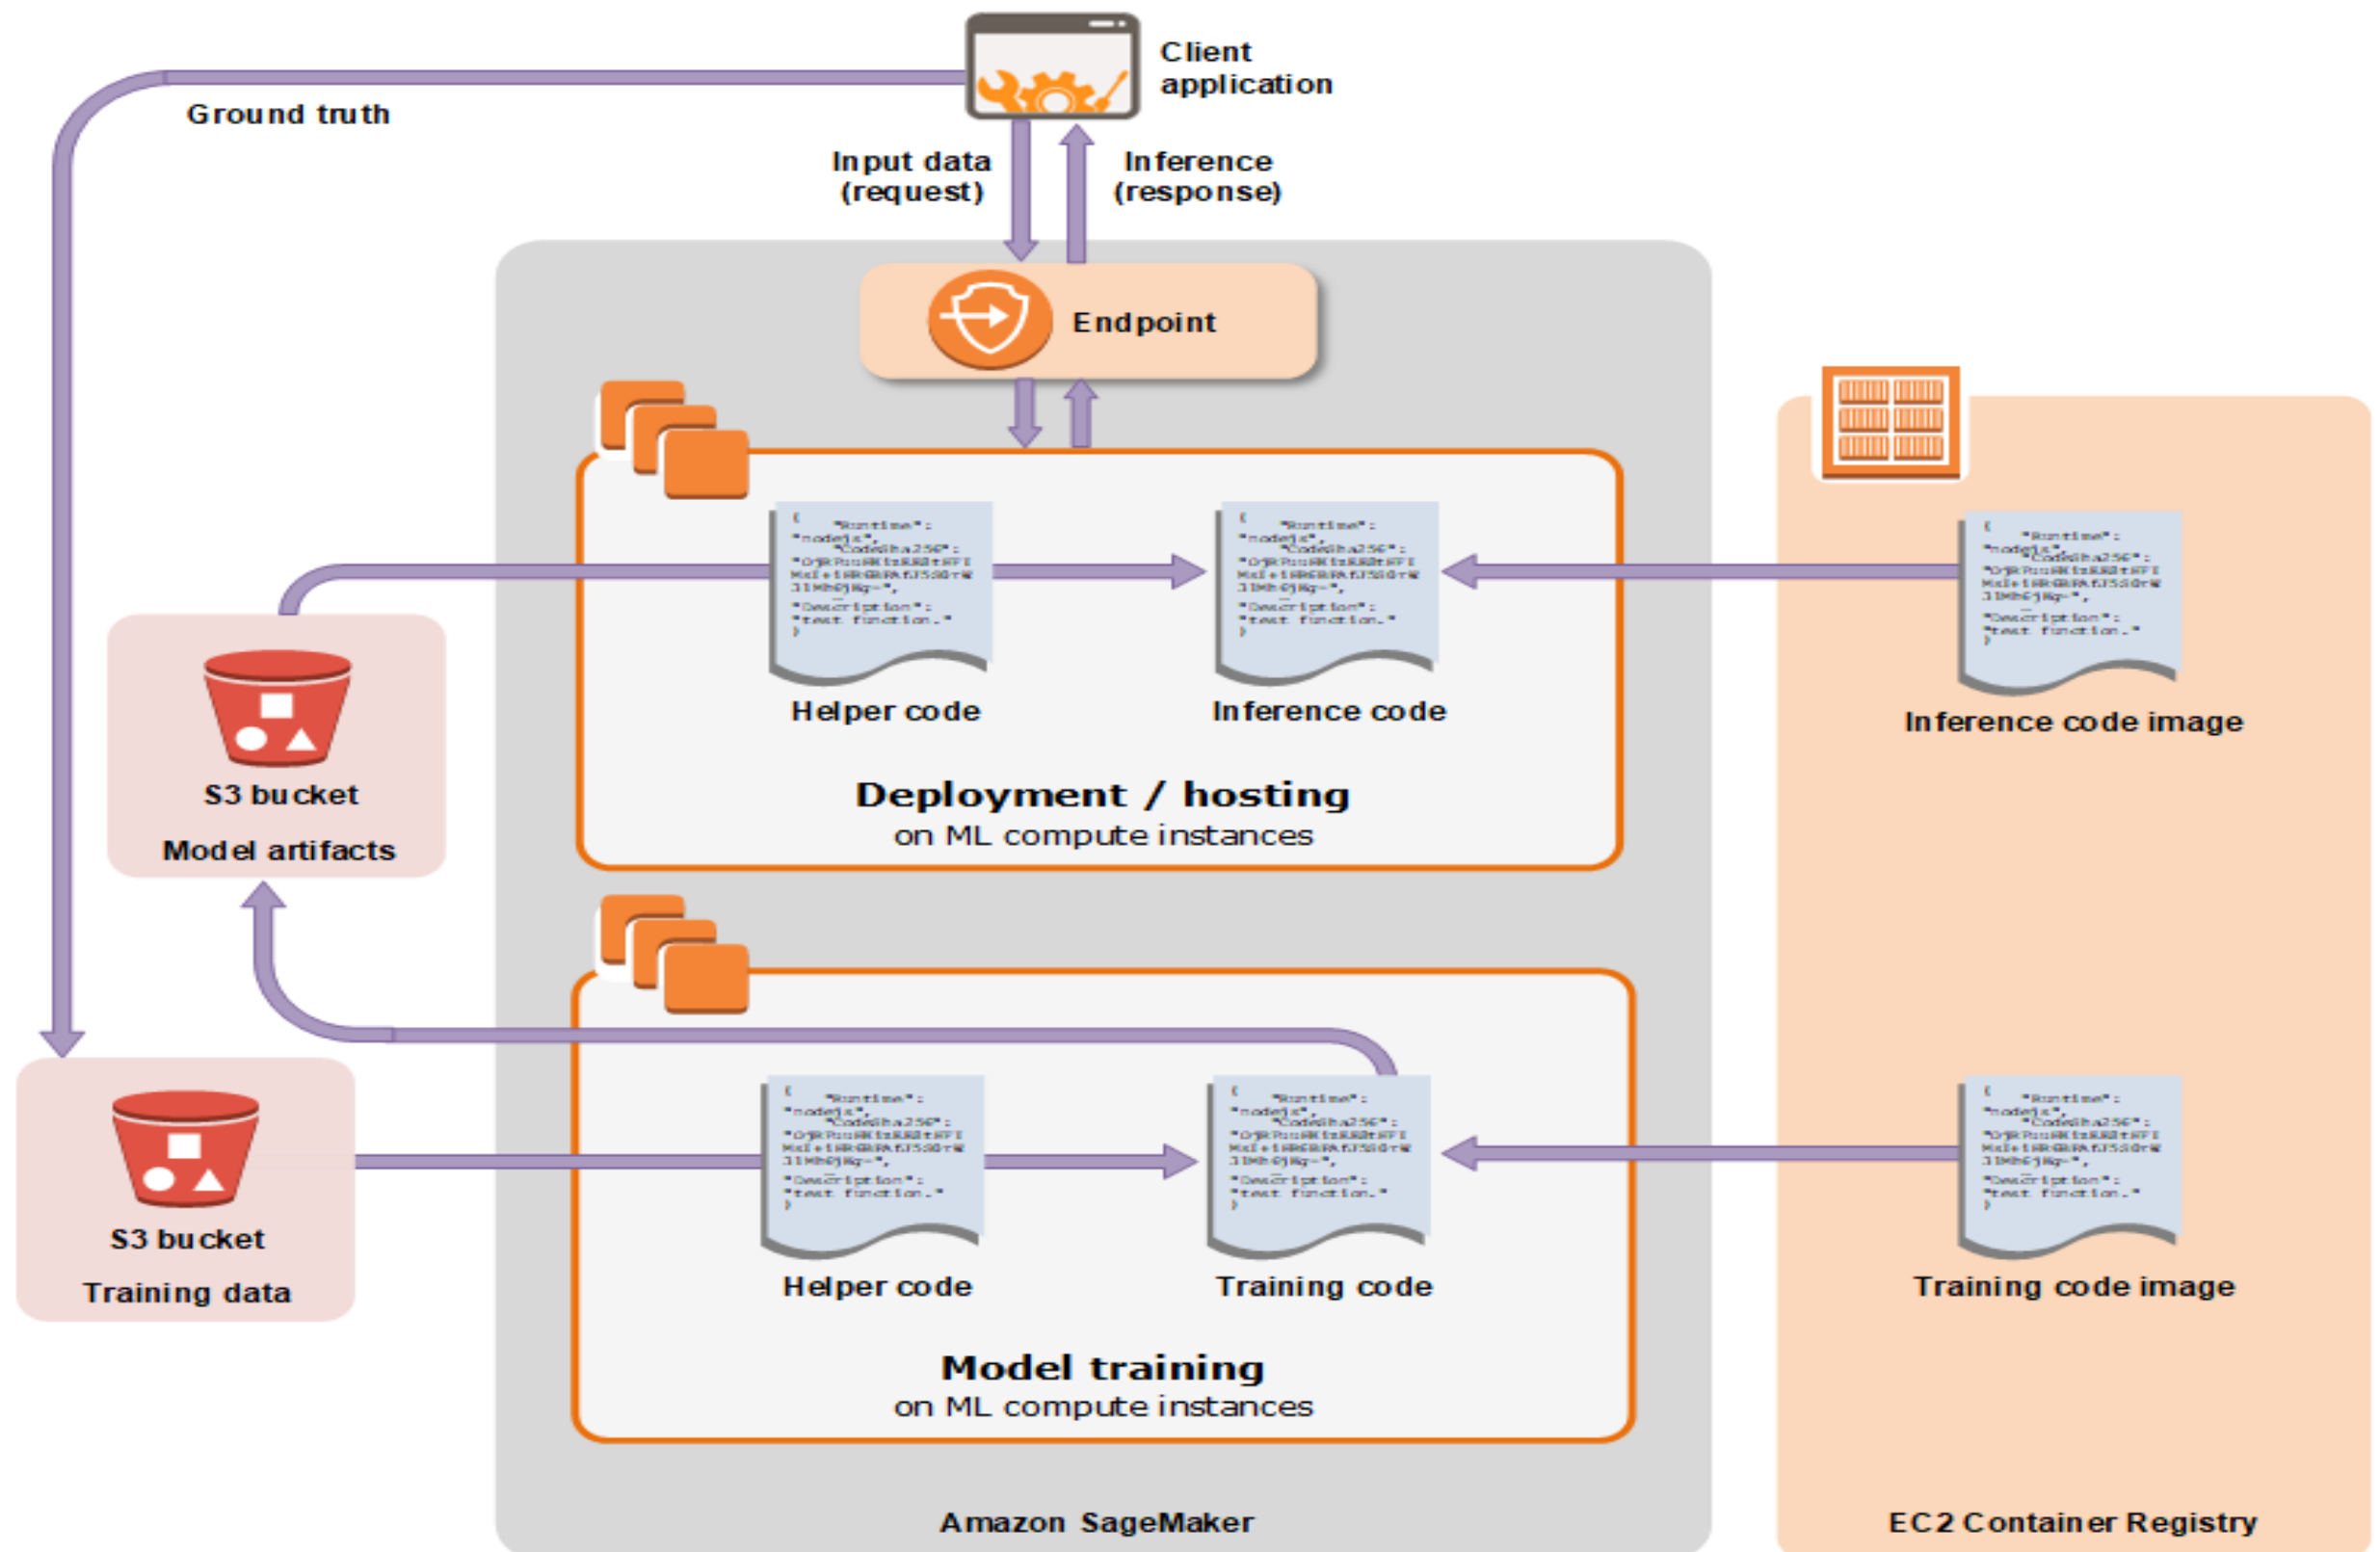
What is Amazon SageMaker?

- Fully managed machine learning service
- Provides an integrated Jupyter authoring notebook instance
- Quickly and easily build, train, and deploy machine learning models
- Bring-your-own-algorithms and frameworks
- Provides common machine learning algorithms
- Secure and scalable environment
- Automatic Model Tuning/Hyperparameter Tuning

Amazon SageMaker - Overview



ML using Amazon SageMaker – Architectural Diagram



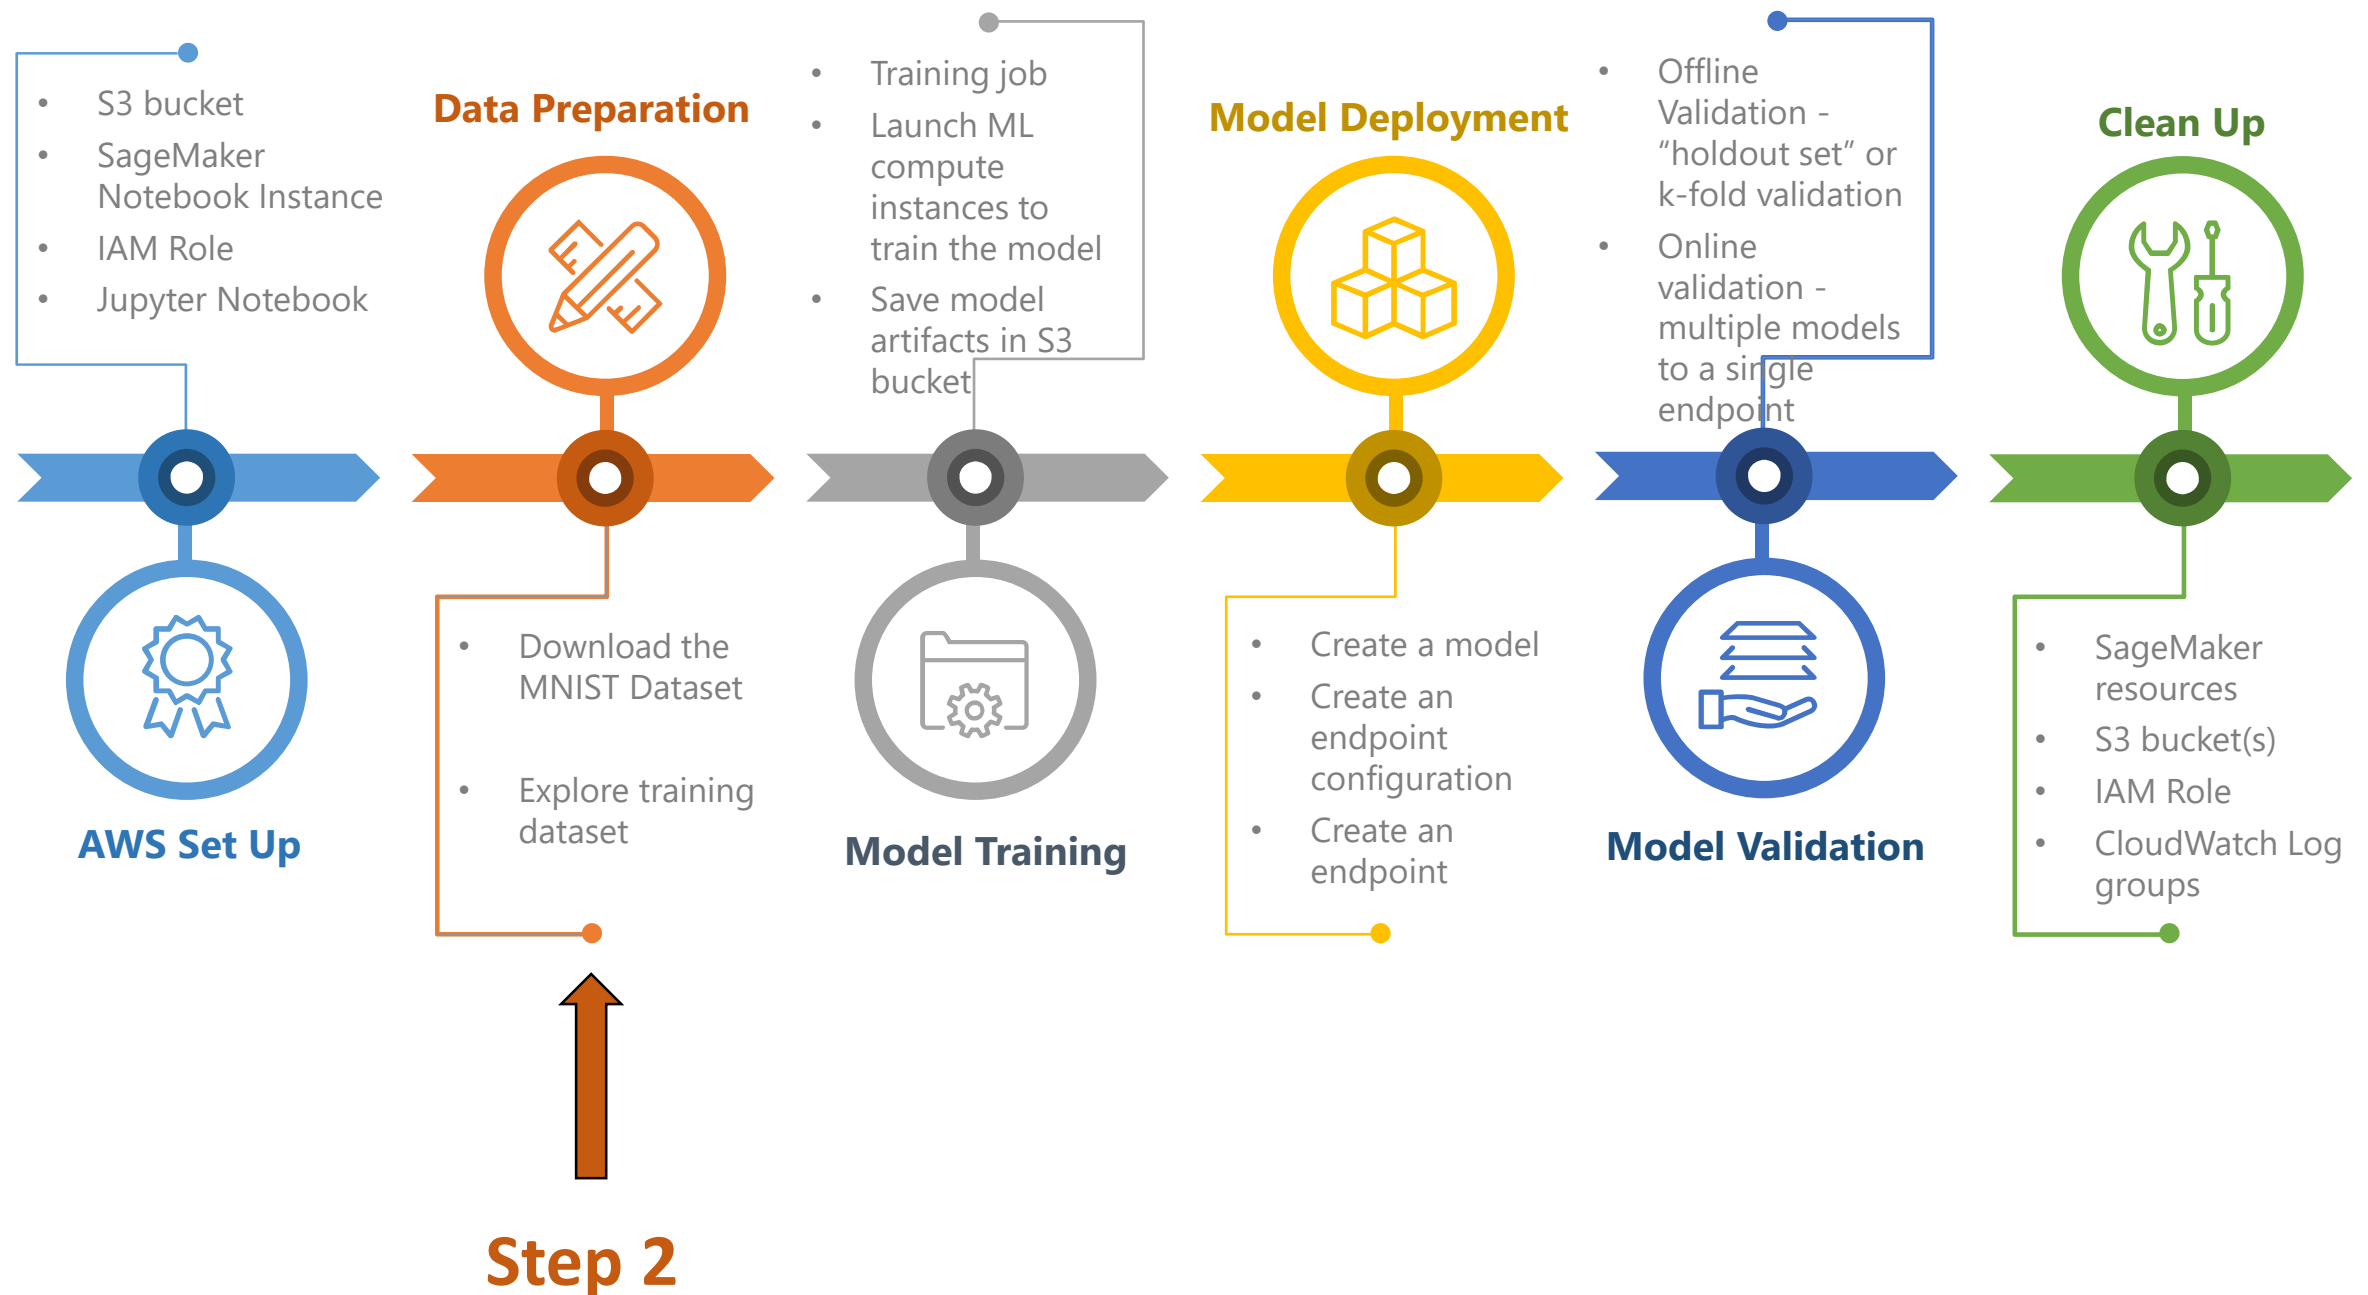
Steps in Machine Learning using Amazon SageMaker



Step 1 – AWS Set up...

- AWS S3 bucket:
 - to save model training data and artifacts
 - Include “sagemaker” in the bucket name
- Amazon SageMaker Notebook Instance:
 - fully managed machine learning (ML) EC2 compute instance running Jupyter Notebook App
- IAM Role (Identity and Access Management):
 - Allows SageMaker to perform actions in other AWS services on your behalf
 - Associate “AmazonSageMakerFullAccess” IAM policy to the IAM role you create
- When notebook instance is “InService”, open the Jupyter dashboard
- Create a Jupyter Notebook
 - Files tab → New → conda_python3 → Name the notebook

Steps in Machine Learning using Amazon SageMaker



Steps in Machine Learning using Amazon SageMaker



Steps in Machine Learning using Amazon SageMaker



Steps in Machine Learning using Amazon SageMaker



Steps in Machine Learning using Amazon SageMaker



Step 6 – Clean up...

To avoid incurring unnecessary charges, use the AWS Console to delete following resources

- Amazon SageMaker console:
 - Endpoint (auto deletes ML compute instances)
 - Endpoint configuration
 - Model
 - Notebook instance (stop the instance before deleting it)
- Amazon S3 console:
 - S3 buckets created for storing model artifacts and training dataset
- IAM console:
 - Delete IAM Role
 - Delete customized IAM Policy, if created
- Amazon CloudWatch console:
 - Log groups starting with “/aws/sagemaker/”

Amazon SageMaker Programming Model

- Use the Amazon SageMaker console
 - Works well for simple jobs
- Modify the example Jupyter notebooks
 - Start with a notebook that has a suitable algorithm
 - Modify it to accommodate your data source and specific needs
- Write model training and inference code from scratch:
 - High-level Python library:
 - Simplifies model training and deployment
 - Handles various operations implicitly
 - AWS SDK:
 - Available in multiple languages and platforms
 - Provides methods that corresponds to SageMaker APIs
 - No need to write authentication code. Uses access keys to implicitly authenticate the requests.
- Integrate Amazon SageMaker into your Apache Spark workflow
 - SageMaker provides a library for calling its APIs from Apache Spark
 - With it, Amazon SageMaker-based estimators can be used in an Apache Spark pipeline

K-Means Algorithm

- Unsupervised classification machine learning algorithm
- Finds discrete groupings within data by training a model that groups similar objects together
- Maps each observation to a point in the n-dimensional space
- Choose the number of clusters (k) to create
- Determines the initial K cluster centers: random or k-means++ approach
- Iterates over the training dataset and calculates cluster centers: moves these centers toward the true cluster centers
- K-Means Hyperparameters
 - **K**: number of required clusters
 - **feature_dim**: number of features in the input data
 - **mini_batch_size**: number of observations per mini-batch
 - **init_method**: random or kmeans++
 - **extra_center_factor**: $K \text{ centers} = \text{num of clusters (k)} * \text{extra_center_factor (x)}$
 - **objective metric**: msd (Mean squared distances)/ssd (Sum of the squared distances)

Automatic Model Tuning/Hyperparameter Tuning

- Supervised machine learning regression problem
- Finds the best version of a model
- Specify the range of hyperparameters values to search to find the best values
- Launches training jobs on your dataset using the algorithm and ranges of hyperparameters
- Chooses the hyperparameter values that result in a model that performs the best, as measured by an objective metric that you choose
- Deploy the best trained model that training job created



Machine Learning using other AWS Services

Demo of Machine Learning using other AWS Services

- **Amazon Comprehend:** Natural Language Processing and Text Analytics
- **Amazon Rekognition:** Deep learning-based visual analysis service for images and videos
- **Amazon Translate:** Natural and fluent translation -automatic detection and translation of a language text
- **Amazon Polly:** Text-to-Speech
- **Amazon Lex:** Build voice and text natural language chatbots
- **AWS DeepLens:** Deep learning-enabled video camera
- **Amazon Transcribe:** Automatic Speech Recognition

Questions ??

Thanks for coming



Have a nice day!